

Decision Tree: Income Prediction

In this case study, we will build a decision tree to predict the income of a given population, which is labelled as $\leq 50K$ and $> 50K$. The attributes (predictors) are age, working class type, marital status, gender, race etc.

In the following sections, we'll:

- clean and prepare the data,
- build a decision tree with default hyperparameters,
- understand all the hyperparameters that we can tune, and finally
- choose the optimal hyperparameters using grid search cross-validation.

Understanding and Cleaning the Data

```
In [1]: # Importing the required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: # To ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: # Reading the csv file and putting it into 'df' object.
df = pd.read_csv('adult_dataset.csv')
```

```
In [4]: # Let's understand the type of values in each column of our dataframe 'df'.  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 32561 entries, 0 to 32560  
Data columns (total 15 columns):  
age                32561 non-null int64  
workclass          32561 non-null object  
fnlwgt             32561 non-null int64  
education          32561 non-null object  
education.num      32561 non-null int64  
marital.status     32561 non-null object  
occupation         32561 non-null object  
relationship       32561 non-null object  
race               32561 non-null object  
sex                32561 non-null object  
capital.gain       32561 non-null int64  
capital.loss       32561 non-null int64  
hours.per.week     32561 non-null int64  
native.country     32561 non-null object  
income             32561 non-null object  
dtypes: int64(6), object(9)  
memory usage: 3.7+ MB
```

```
In [5]: # Let's understand the data, how it look like.
df.head(20)
```

```
Out[5]:
```

	class	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	cap
	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	
	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	
	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	
	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female	
	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male	
	Self-gov	88638	Doctorate	16	Never-married	Prof-specialty	Other-relative	White	Female	
	Local-gov	422013	HS-grad	9	Divorced	Prof-specialty	Not-in-family	White	Female	
	Private	70037	Some-college	10	Never-married	Craft-repair	Unmarried	White	Male	
	Private	172274	Doctorate	16	Divorced	Prof-specialty	Unmarried	Black	Female	
	Unemp-not-inc	164526	Prof-school	15	Never-married	Prof-specialty	Not-in-family	White	Male	
	Private	129177	Bachelors	13	Widowed	Other-service	Not-in-family	White	Female	
	Private	136204	Masters	14	Separated	Exec-managerial	Not-in-family	White	Male	
	?	172175	Doctorate	16	Never-married	?	Not-in-family	White	Male	
	Private	45363	Prof-school	15	Divorced	Prof-specialty	Not-in-family	White	Male	
	Private	172822	11th	7	Divorced	Transport-moving	Not-in-family	White	Male	
	Private	317847	Masters	14	Divorced	Exec-managerial	Not-in-family	White	Male	
	Private	119592	Assoc-acdm	12	Never-married	Handlers-cleaners	Not-in-family	Black	Male	
	Private	203034	Bachelors	13	Separated	Sales	Not-in-family	White	Male	

You can observe that the columns workclass and occupation consist of missing values which are represented as '?' in the dataframe.

```
In [6]: # rows with missing values represented as '?'.
df_1 = df[df.workclass == '?']
df_1
```

```
Out[6]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family
2	66	?	186061	Some-college	10	Widowed	?	Unmarried
14	51	?	172175	Doctorate	16	Never-married	?	Not-in-family
24	61	?	135285	HS-grad	9	Married-civ-spouse	?	Husband
44	71	?	100820	HS-grad	9	Married-civ-spouse	?	Husband
48	68	?	192052	Some-college	10	Married-civ-spouse	?	Wife
49	67	?	174995	Some-college	10	Married-civ-spouse	?	Husband
76	41	?	27187	Assoc-voc	11	Married-civ-spouse	?	Husband

Now we can check the number of rows in df_1.

```
In [7]: df_1.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1836 entries, 0 to 32544
Data columns (total 15 columns):
age                1836 non-null int64
workclass          1836 non-null object
fnlwgt             1836 non-null int64
education          1836 non-null object
education.num      1836 non-null int64
marital.status     1836 non-null object
occupation         1836 non-null object
relationship       1836 non-null object
race              1836 non-null object
sex               1836 non-null object
capital.gain       1836 non-null int64
capital.loss       1836 non-null int64
hours.per.week     1836 non-null int64
native.country     1836 non-null object
income            1836 non-null object
dtypes: int64(6), object(9)
memory usage: 229.5+ KB
```

There are 1836 rows with missing values, which is about 5% of the total data. We choose to simply drop these rows.

```
In [8]: # dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()
```

```
Out[8]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	rac
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	Whit
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	Whit
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	Whit
5	34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	Whit
6	38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	Whit

Let's see whether any other columns contain a "?". Since "?" is a string, we can apply this check only on the categorical columns.

```
In [9]: # select all categorical variables
df_categorical = df.select_dtypes(include=['object'])

# checking whether any other columns contain a "?"
df_categorical.apply(lambda x: x=="?", axis=0).sum()
```

```
Out[9]: workclass      0
education    0
marital.status 0
occupation    7
relationship  0
race          0
sex           0
native.country 556
income        0
dtype: int64
```

Thus, the columns occupation and native.country contain some "?"s. Let's get rid of them.

```
In [10]: # dropping the "?"s
df = df[df['occupation'] != '?']
df = df[df['native.country'] != '?']
```

Now we have a clean dataframe which is ready for model building.

```
In [11]: # clean dataframe
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
age                30162 non-null int64
workclass          30162 non-null object
fnlwgt            30162 non-null int64
education          30162 non-null object
education.num      30162 non-null int64
marital.status     30162 non-null object
occupation        30162 non-null object
relationship       30162 non-null object
race              30162 non-null object
sex               30162 non-null object
capital.gain       30162 non-null int64
capital.loss       30162 non-null int64
hours.per.week     30162 non-null int64
native.country     30162 non-null object
income            30162 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

Data Preparation

There are a number of preprocessing steps we need to do before building the model.

Firstly, note that we have both categorical and numeric features as predictors. In previous models such as linear and logistic regression, we had created **dummy variables** for categorical variables, since those models (being mathematical equations) can process only numeric variables.

All that is not required in decision trees, since they can process categorical variables easily. However, we still need to **encode the categorical variables** into a standard format so that sklearn can understand them and build the tree. We'll do that using the `LabelEncoder()` class, which comes with `sklearn.preprocessing`.

```
In [12]: from sklearn import preprocessing

# encode categorical variables using Label Encoder

# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()
```

```
Out[12]:
```

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	inc
1	Private	HS-grad	Widowed	Exec-managerial	Not-in-family	White	Female	United-States	<
3	Private	7th-8th	Divorced	Machine-op-inspct	Unmarried	White	Female	United-States	<
4	Private	Some-college	Separated	Prof-specialty	Own-child	White	Female	United-States	<
5	Private	HS-grad	Divorced	Other-service	Unmarried	White	Female	United-States	<
6	Private	10th	Separated	Adm-clerical	Unmarried	White	Male	United-States	<

```
In [13]: # apply Label encoder to df_categorical

le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

```
Out[13]:
```

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
1	2	11	6	3	1	4	0	38	0
3	2	5	0	6	4	4	0	38	0
4	2	15	5	9	3	4	0	38	0
5	2	11	0	7	4	4	0	38	0
6	2	0	5	0	4	4	1	38	0

```
In [14]: # concat df_categorical with original df
df = df.drop(df_categorical.columns, axis=1)
df = pd.concat([df, df_categorical], axis=1)
df.head()
```

```
Out[14]:
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	marital.status
1	82	132870	9	0	4356	18	2	11	Married
3	54	140359	4	0	3900	40	2	5	Married
4	41	264663	10	0	3900	40	2	15	Married
5	34	216864	9	0	3770	45	2	11	Married
6	38	150601	6	0	3770	40	2	0	Married

```
In [15]: # Look at column types
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 1 to 32560
Data columns (total 15 columns):
age                30162 non-null int64
fnlwgt             30162 non-null int64
education.num      30162 non-null int64
capital.gain       30162 non-null int64
capital.loss       30162 non-null int64
hours.per.week     30162 non-null int64
workclass          30162 non-null int32
education          30162 non-null int32
marital.status     30162 non-null int32
occupation         30162 non-null int32
relationship       30162 non-null int32
race               30162 non-null int32
sex               30162 non-null int32
native.country     30162 non-null int32
income             30162 non-null int32
dtypes: int32(9), int64(6)
memory usage: 2.6 MB
```

```
In [16]: # convert target variable income to categorical - Target variable not needed in
Decision node to Check
#df['income'] = df['income'].astype('category')
#df.head()
```

Now all the categorical variables are suitably encoded. Let's build the model.

Model Building and Evaluation

Let's first build a decision tree with default hyperparameters. Then we'll use cross-validation to tune them.

```
In [17]: # Importing train-test-split
         from sklearn.model_selection import train_test_split
```

```
In [18]: # Putting feature variable to X
X = df.drop('income',axis=1)

# Putting response variable to y
y = df['income']
```

[illegible]

```
Out[19]:
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	workclass	education	mari
1	42	289636	9	0	0	46	2	11	
6	37	52465	9	0	0	40	1	11	
7	38	125933	14	0	0	40	0	12	
2	44	183829	13	0	0	38	5	9	
3	35	198841	11	0	0	35	2	8	

```
In [20]: # Importing decision tree classifier from sklearn library
          from sklearn.tree import DecisionTreeClassifier

          # Fitting the decision tree with default hyperparameters, apart from
          # max_depth which is 5 so that we can plot and read the tree.
          dt_default = DecisionTreeClassifier(max_depth=5)
          dt_default.fit(X_train, y_train)
```

[illegible]

```
In [21]: # Let's check the evaluation metrics of our default model

# Importing classification report and confusion matrix from sklearn metrics
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score

# Making predictions
y_pred_default = dt_default.predict(X_test)

# Printing classification report
print(classification_report(y_test, y_pred_default))
```

	precision	recall	f1-score	support
0	0.86	0.95	0.91	6867
1	0.78	0.52	0.63	2182
accuracy			0.85	9049
macro avg	0.82	0.74	0.77	9049
weighted avg	0.84	0.85	0.84	9049

```
In [22]: # Printing confusion matrix and accuracy
print(confusion_matrix(y_test,y_pred_default))
print(accuracy_score(y_test,y_pred_default))
```

```
[[6553  314]
 [1038 1144]]
0.8505912255497845
```

Hyperparameter Tuning

Tuning max_depth

Let's first try to find the optimum values for max_depth and understand how the value of max_depth affects the decision tree.

Here, we are creating a dataframe with max_depth in range 1 to 80 and checking the accuracy score corresponding to each max_depth.

To reiterate, a grid search scheme consists of:

- an estimator (classifier such as SVC() or decision tree)
- a parameter space
- a method for searching or sampling candidates (optional)
- a cross-validation scheme, and
- a score function (accuracy, roc_auc etc.)

```
In [55]: # GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

# specify number of folds for k-fold CV
n_folds = 5

# parameters to build the model on
parameters = {'max_depth': range(1, 40)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)
```

```
Out[55]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=DecisionTreeClassifier(class_weight=None, criterion='gini', ma
x_depth=None,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort=False, random_state=100,
                      splitter='best'),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'max_depth': range(1, 40)}, pre_dispatch='2*n_jobs',
                      refit=True, return_train_score='warn', scoring='accuracy',
                      verbose=0)
```

```
In [56]: # scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```


```
Out[56]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	params	score
0	0.013414	0.006326	0.001799	0.000979	1	{'max_depth': 1}	0.825
1	0.018751	0.006248	0.006253	0.007658	2	{'max_depth': 2}	0.845
2	0.018158	0.003160	0.006247	0.007651	3	{'max_depth': 3}	0.855
3	0.051170	0.018552	0.002007	0.004013	4	{'max_depth': 4}	0.855
4	0.035856	0.007133	0.006250	0.007654	5	{'max_depth': 5}	0.855

5 rows × 21 columns

Now let's visualize how train and test score changes with max_depth.

```
In [57]: # plotting accuracies with max_depth
plt.figure()
plt.plot(scores["param_max_depth"],
         scores["mean_train_score"],
         label="training accuracy")
plt.plot(scores["param_max_depth"],
         scores["mean_test_score"],
         label="test accuracy")
plt.xlabel("max_depth")
plt.ylabel("Accuracy")
plt.legend()
```



You can see that as we increase the value of `max_depth`, both training and test score increase till about `max_depth = 10`, after which the test score gradually reduces. Note that the scores are average accuracies across the 5-folds.

Thus, it is clear that the model is overfitting the training data if the `max_depth` is too high. Next, let's see how the model behaves with other hyperparameters.

Tuning `min_samples_leaf`

The hyperparameter **`min_samples_leaf`** indicates the minimum number of samples required to be at a leaf.

So if the values of `min_samples_leaf` is less, say 5, then the will be constructed even if a leaf has 5, 6 etc. observations (and is likely to overfit).

Let's see what will be the optimum value for `min_samples_leaf`.

```
In [39]: # GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

# specify number of folds for k-fold CV
n_folds = 5

# parameters to build the model on
parameters = {'min_samples_leaf': range(5, 200, 20)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)
```


```
Out[39]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=DecisionTreeClassifier(class_weight=None, criterion='gini', ma
x_depth=None,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort=False, random_state=100,
                      splitter='best'),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'min_samples_leaf': range(5, 200, 20)},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring='accuracy', verbose=0)
```

```
In [32]: # scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```


```
Out[32]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_min_samples_leaf
0	0.076940	0.002072	0.823663	0.913785	5 {
1	0.070042	0.002049	0.844172	0.869180	25 {
2	0.055783	0.001711	0.848529	0.861554	45 {
3	0.061370	0.002087	0.846114	0.857067	65 {
4	0.051622	0.002022	0.845451	0.854320	85 {

5 rows × 21 columns



```
In [33]: # plotting accuracies with min_samples_leaf
plt.figure()
plt.plot(scores["param_min_samples_leaf"],
         scores["mean_train_score"],
         label="training accuracy")
plt.plot(scores["param_min_samples_leaf"],
         scores["mean_test_score"],
         label="test accuracy")
plt.xlabel("min_samples_leaf")
plt.ylabel("Accuracy")
plt.legend()
```



You can see that at low values of `min_samples_leaf`, the tree gets a bit overfitted. At values > 100 , however, the model becomes more stable and the training and test accuracy start to converge.

Tuning `min_samples_split`

The hyperparameter `min_samples_split` is the minimum no. of samples required to split an internal node. Its default value is 2, which means that even if a node is having 2 samples it can be further divided into leaf nodes.

```
In [40]: # GridSearchCV to find optimal min_samples_split
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

# specify number of folds for k-fold CV
n_folds = 5

# parameters to build the model on
parameters = {'min_samples_split': range(5, 200, 20)}

# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
                              random_state = 100)

# fit tree on training data
tree = GridSearchCV(dtree, parameters,
                    cv=n_folds,
                    scoring="accuracy")
tree.fit(X_train, y_train)
```

```
Out[40]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=DecisionTreeClassifier(class_weight=None, criterion='gini', ma
x_depth=None,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort=False, random_state=100,
                      splitter='best'),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'min_samples_split': range(5, 200, 20)},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring='accuracy', verbose=0)
```


```
In [41]: # scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

```
Out[41]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_min_samples_split	
0	0.249302	0.023308	0.007176	0.001296	5	{'min_sa
1	0.254080	0.041291	0.006585	0.002792	25	{'min_sa
2	0.248087	0.022977	0.008883	0.002362	45	{'min_sa
3	0.271269	0.056348	0.008110	0.001872	65	{'min_sa
4	0.216440	0.032674	0.007011	0.002618	85	{'min_sa

5 rows × 21 columns


```
In [36]: # plotting accuracies with min_samples_leaf
plt.figure()
plt.plot(scores["param_min_samples_split"],
         scores["mean_train_score"],
         label="training accuracy")
plt.plot(scores["param_min_samples_split"],
         scores["mean_test_score"],
         label="test accuracy")
plt.xlabel("min_samples_split")
plt.ylabel("Accuracy")
plt.legend()
```



This shows that as you increase the `min_samples_split`, the tree overfits lesser since the model is less complex.

Grid Search to Find Optimal Hyperparameters

We can now use `GridSearchCV` to find multiple optimal hyperparameters together. Note that this time, we'll also specify the criterion (gini/entropy or IG).

```
In [58]: # Create the parameter grid
param_grid = {
    'max_depth': range(5, 15, 5),
    'min_samples_leaf': range(50, 150, 50),
    'min_samples_split': range(50, 150, 50),
    'criterion': ["entropy", "gini"]
}

n_folds = 5

# Instantiate the grid search model
dtree = DecisionTreeClassifier()
grid_search = GridSearchCV(estimator = dtree, param_grid = param_grid,
                           cv = n_folds, verbose = 1)

# Fit the grid search to the data
grid_search.fit(X_train,y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent worker
s.
[Parallel(n_jobs=1)]: Done 80 out of 80 | elapsed: 4.4s finished
```

```
Out[58]: GridSearchCV(cv=5, error_score='raise-deprecating',
    estimator=DecisionTreeClassifier(class_weight=None, criterion='gini', ma
x_depth=None,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, presort=False, random_state=None,
    splitter='best'),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid={'max_depth': range(5, 15, 5), 'min_samples_leaf': range(50,
150, 50), 'min_samples_split': range(50, 150, 50), 'criterion': ['entropy', 'gi
ni']},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring=None, verbose=1)
```

```
In [59]: # cv results
cv_results = pd.DataFrame(grid_search.cv_results_)
cv_results
```

```
Out[59]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_criterion	param_max_de
0	0.050325	0.008642	0.001601	0.001358	entropy	
1	0.032814	0.003100	0.003125	0.006250	entropy	
2	0.038284	0.004984	0.003126	0.006251	entropy	
3	0.040109	0.006026	0.000000	0.000000	entropy	

```
In [60]: # printing the optimal accuracy score and hyperparameters
print("best accuracy", grid_search.best_score_)
print(grid_search.best_estimator_)
```

```
best accuracy 0.8514659214701843
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=10,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=50, min_samples_split=50,
                        min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                        splitter='best')
```

Running the model with best parameters obtained from grid search.

```
In [61]: # model with optimal hyperparameters
clf_gini = DecisionTreeClassifier(criterion = "gini",
                                random_state = 100,
                                max_depth=10,
                                min_samples_leaf=50,
                                min_samples_split=50)

clf_gini.fit(X_train, y_train)
```

```
Out[61]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=10,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=50, min_samples_split=50,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=100,
                                splitter='best')
```

```
In [62]: # accuracy score
clf_gini.score(X_test,y_test)
```

```
Out[62]: 0.850922753895458
```

```
In [63]: # plotting the tree
dot_data = StringIO()
export_graphviz(clf_gini,
                out_file=dot_data,feature_names=features,filled=True,rounded=True)

graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-63-62884116144f> in <module>()
      1 # plotting the tree
      2 dot_data = StringIO()
----> 3 export_graphviz(clf_gini, out_file=dot_data,feature_names=features,fill
      ed=True,rounded=True)
      4
      5 graph = pydotplus.graph_from_dot_data(dot_data.getvalue())

NameError: name 'features' is not defined
```

You can see that this tree is too complex to understand. Let's try reducing the max_depth and see how the tree looks.

```
In [64]: # tree with max_depth = 3
clf_gini = DecisionTreeClassifier(criterion = "gini",
                                random_state = 100,
                                max_depth=3,
                                min_samples_leaf=50,
                                min_samples_split=50)

clf_gini.fit(X_train, y_train)

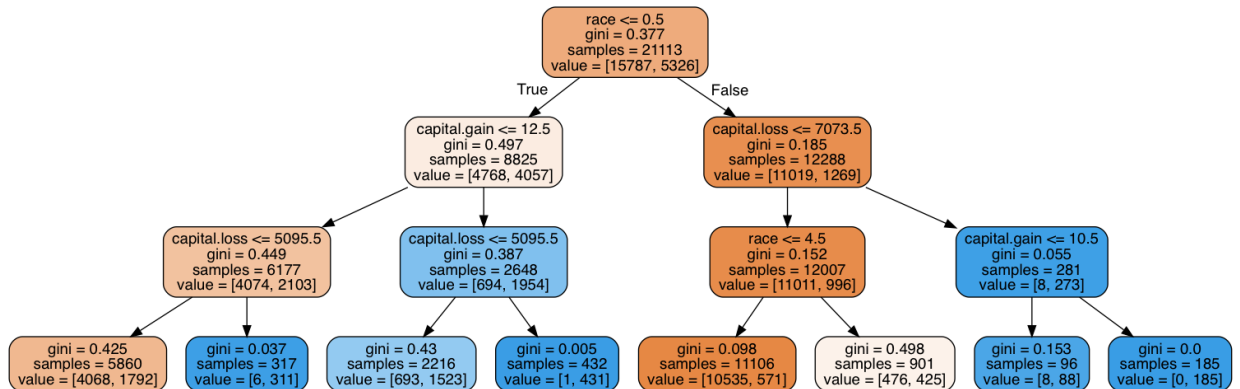
# score
print(clf_gini.score(X_test,y_test))

0.8393192617968837
```

```
In [44]: # plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(clf_gini,
out_file=dot_data,feature_names=features,filled=True,rounded=True)

graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```

Out[44]:



```
In [45]: # classification metrics
from sklearn.metrics import classification_report, confusion_matrix
y_pred = clf_gini.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	6867
1	0.77	0.47	0.59	2182
avg / total	0.83	0.84	0.82	9049

```
In [46]: # confusion matrix
print(confusion_matrix(y_test,y_pred))

[[6564  303]
 [1151 1031]]
```

