

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
sns.set_theme(color_codes=True)
```

```
In [2]: # import dataset
df = pd.read_csv('realtor-data.csv')
df.head()
```

Out[2]:

	status	bed	bath	acre_lot	city	state	zip_code	house_size	prev_sold_date	price
0	for_sale	3.0	2.0	0.12	Adjuntas	Puerto Rico	601.0	920.0	NaN	105000.0
1	for_sale	4.0	2.0	0.08	Adjuntas	Puerto Rico	601.0	1527.0	NaN	80000.0
2	for_sale	2.0	1.0	0.15	Juana Diaz	Puerto Rico	795.0	748.0	NaN	67000.0
3	for_sale	4.0	2.0	0.10	Ponce	Puerto Rico	731.0	1800.0	NaN	145000.0
4	for_sale	6.0	2.0	0.05	Mayaguez	Puerto Rico	680.0	NaN	NaN	65000.0

## Data Preprocessing Part 1

```
In [3]: #Check the number of unique value for object datatypes
df.select_dtypes(include='object').nunique()
```

```
Out[3]: status          2
city          525
state         12
prev_sold_date 3604
dtype: int64
```

```
In [4]: #Check the number of missing value
df.prev_sold_date.isnull().sum()
```

Out[4]: 71255

```
In [5]: #Check the number of row and column
df.shape
```

Out[5]: (100000, 10)

```
In [6]: # Drop prev_sold_date because the missing value is around 70%
df.drop(columns='prev_sold_date', inplace=True)
df.head()
```

Out[6]:

	status	bed	bath	acre_lot	city	state	zip_code	house_size	price
0	for_sale	3.0	2.0	0.12	Adjuntas	Puerto Rico	601.0	920.0	105000.0
1	for_sale	4.0	2.0	0.08	Adjuntas	Puerto Rico	601.0	1527.0	80000.0
2	for_sale	2.0	1.0	0.15	Juana Diaz	Puerto Rico	795.0	748.0	67000.0
3	for_sale	4.0	2.0	0.10	Ponce	Puerto Rico	731.0	1800.0	145000.0
4	for_sale	6.0	2.0	0.05	Mayaguez	Puerto Rico	680.0	NaN	65000.0

```
In [7]: # Drop city because the number of unique value for object datatypes is alot
df.drop(columns='city', inplace=True)
df.head()
```

Out[7]:

	status	bed	bath	acre_lot	state	zip_code	house_size	price
0	for_sale	3.0	2.0	0.12	Puerto Rico	601.0	920.0	105000.0
1	for_sale	4.0	2.0	0.08	Puerto Rico	601.0	1527.0	80000.0
2	for_sale	2.0	1.0	0.15	Puerto Rico	795.0	748.0	67000.0
3	for_sale	4.0	2.0	0.10	Puerto Rico	731.0	1800.0	145000.0
4	for_sale	6.0	2.0	0.05	Puerto Rico	680.0	NaN	65000.0

## Exploratory Data Analysis


```
In [8]: # List of categorical variables to plot
cat_vars = ['status', 'state']

# create figure with subplots
fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(20, 10))
axs = axs.flatten()

# create barplot for each categorical variable
for i, var in enumerate(cat_vars):
    sns.barplot(x=var, y='price', data=df, ax=axs[i])
    axs[i].set_xticklabels(axs[i].get_xticklabels(), rotation=90)

# adjust spacing between subplots
fig.tight_layout()

# show plot
plt.show()
```




```
In [9]: num_vars = ['bed', 'bath', 'acre_lot', 'house_size']

fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.boxplot(x=var, data=df, ax=axs[i])

fig.tight_layout()

plt.show()
```



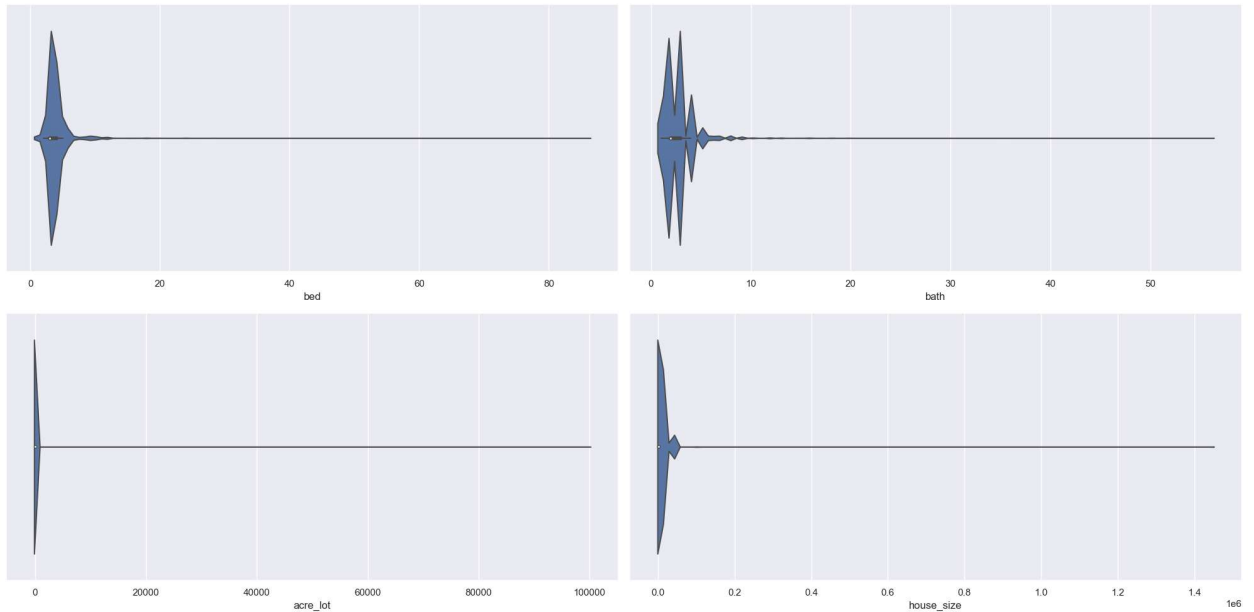
```
In [10]: num_vars = ['bed', 'bath', 'acre_lot', 'house_size']

fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.violinplot(x=var, data=df, ax=axs[i])

fig.tight_layout()

plt.show()
```




```
In [11]: num_vars = ['bed', 'bath', 'acre_lot', 'house_size']

fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.scatterplot(x=var, y='price', data=df, ax=axs[i])

fig.tight_layout()

plt.show()
```



## Data Preprocessing Part 2

```
In [12]: check_missing = df.isnull().sum() * 100 / df.shape[0]
check_missing[check_missing > 0].sort_values(ascending=False)
```

```
Out[12]: bed          24.950
house_size  24.918
bath        24.888
acre_lot    14.013
zip_code     0.195
dtype: float64
```

```
In [13]: df.drop(columns='zip_code', inplace=True)
df.shape
```

```
Out[13]: (100000, 7)
```

```
In [14]: # Fill null value with median and mean
df['bed'].fillna(df['bed'].median(), inplace=True)
df['bath'].fillna(df['bath'].median(), inplace=True)
df['house_size'].fillna(df['house_size'].mean(), inplace=True)
df['acre_lot'].fillna(df['acre_lot'].mean(), inplace=True)
```

In [15]: `df.head()`

Out[15]:

	status	bed	bath	acre_lot	state	house_size	price
0	for_sale	3.0	2.0	0.12	Puerto Rico	920.000000	105000.0
1	for_sale	4.0	2.0	0.08	Puerto Rico	1527.000000	80000.0
2	for_sale	2.0	1.0	0.15	Puerto Rico	748.000000	67000.0
3	for_sale	4.0	2.0	0.10	Puerto Rico	1800.000000	145000.0
4	for_sale	6.0	2.0	0.05	Puerto Rico	2180.081737	65000.0

## Label encoding for object datatypes

```
In [16]: # Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:

    # Print the column name and the unique values
    print(f"{col}: {df[col].unique()}")
```

```
status: ['for_sale' 'ready_to_build']
state: ['Puerto Rico' 'Virgin Islands' 'Massachusetts' 'Connecticut'
       'New Hampshire' 'Vermont' 'New Jersey' 'New York' 'South Carolina'
       'Tennessee' 'Rhode Island' 'Virginia']
```

In [17]: `from sklearn import preprocessing`

```
# Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:

    # Initialize a LabelEncoder object
    label_encoder = preprocessing.LabelEncoder()

    # Fit the encoder to the unique values in the column
    label_encoder.fit(df[col].unique())

    # Transform the column using the encoder
    df[col] = label_encoder.transform(df[col])

    # Print the column name and the unique encoded values
    print(f"{col}: {df[col].unique()}")
```

```
status: [0 1]
state: [ 5 10  1  0  2  9  3  4  7  8  6 11]
```

## Remove Outliers using IQR

```
In [18]: # Print the number of rows before outliers removal
df.shape
```

Out[18]: (100000, 7)

```

In [19]: # define a function to remove outliers using IQR
def remove_outliers_iqr(df, columns):
    for col in columns:
        q1 = df[col].quantile(0.25)
        q3 = df[col].quantile(0.75)
        iqr = q3 - q1
        lower_bound = q1 - 1.5 * iqr
        upper_bound = q3 + 1.5 * iqr
        df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]
    return df

# specify the columns to remove outliers from
columns_to_check = ['bed', 'bath', 'acre_lot', 'house_size']

# call the function to remove outliers using IQR
df_clean = remove_outliers_iqr(df, columns_to_check)

# print the resulting dataframe
df_clean.head()

```

Out[19]:

	status	bed	bath	acre_lot	state	house_size	price
0	0	3.0	2.0	0.12	5	920.0	105000.0
1	0	4.0	2.0	0.08	5	1527.0	80000.0
3	0	4.0	2.0	0.10	5	1800.0	145000.0
7	0	3.0	2.0	0.08	5	1050.0	71600.0
10	0	3.0	2.0	13.39	5	1106.0	89000.0

```
In [20]: df_clean.shape
```

Out[20]: (50558, 7)

## Correlation Heatmap

```
In [21]: #Correlation Heatmap
plt.figure(figsize=(20, 16))
sns.heatmap(df_clean.corr(method='pearson'), fmt='.2g', annot=True)
```

Out[21]: <AxesSubplot:>



```
In [23]: df_clean.drop(columns='bath', inplace=True)
df_clean.head()
```

Out[23]:

	status	bed	acre_lot	state	house_size	price
0	0	3.0	0.12	5	920.0	105000.0
1	0	4.0	0.08	5	1527.0	80000.0
3	0	4.0	0.10	5	1800.0	145000.0
7	0	3.0	0.08	5	1050.0	71600.0
10	0	3.0	13.39	5	1106.0	89000.0

## Train Test Split



```
In [24]: X = df_clean.drop('price', axis=1)
y = df_clean['price']
```

```
In [25]: #test size 20% and train size 80%
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,random_state=0)
```

## Decision Tree Regressor

```
In [26]: from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV

# Create a DecisionTreeRegressor object
dtree = DecisionTreeRegressor()

# Define the hyperparameters to tune and their values
param_grid = {
    'max_depth': [2, 4, 6, 8],
    'min_samples_split': [2, 4, 6, 8],
    'min_samples_leaf': [1, 2, 3, 4],
    'max_features': ['auto', 'sqrt', 'log2'],
    'random_state': [0, 42]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(dtree, param_grid, cv=5, scoring='neg_mean_squared_error')

# Fit the GridSearchCV object to the data
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print(grid_search.best_params_)

{'max_depth': 8, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 2,
 'random_state': 42}
```

```
In [27]: from sklearn.tree import DecisionTreeRegressor
dtree = DecisionTreeRegressor(random_state=42, max_depth=8, max_features='auto', min_sam
dtree.fit(X_train, y_train)
```


```
Out[27]: DecisionTreeRegressor(max_depth=8, max_features='auto', random_state=42)
```

```
In [28]: from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = dtree.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

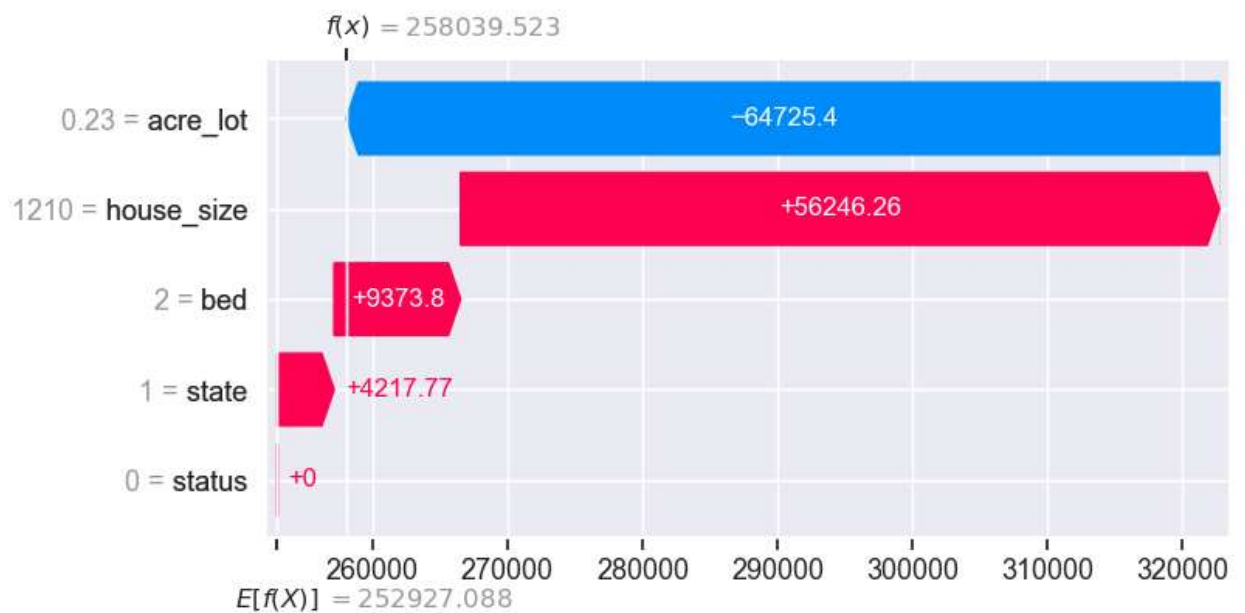
print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))
```

```
MAE is 99905.67946085877
MAPE is 1.3304384099501423
MSE is 50681781752.41356
R2 score is 0.5766792701837287
RMSE score is 225126.14631004893
```

```
In [29]: import shap
explainer = shap.TreeExplainer(dtree)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```

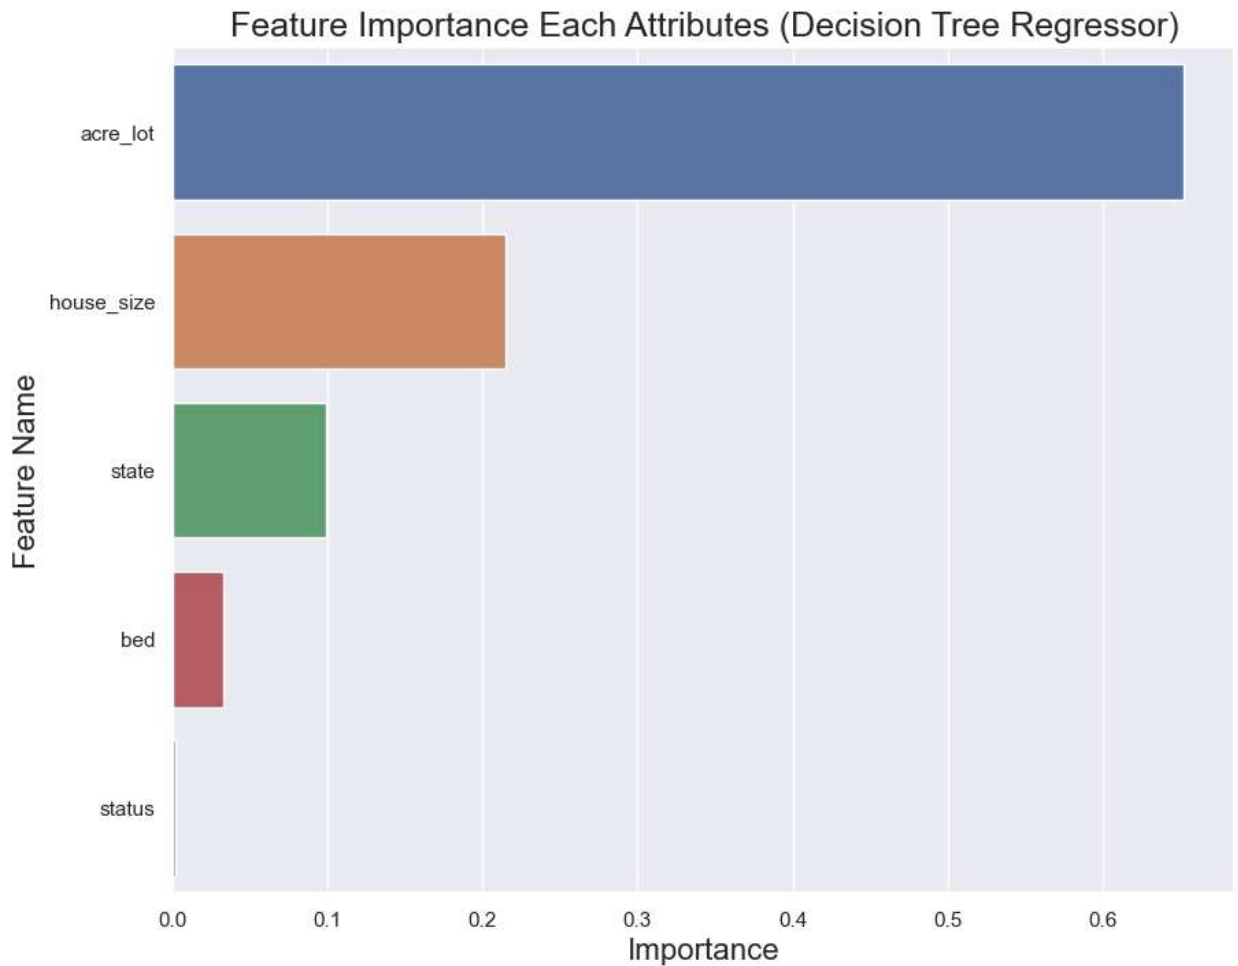


```
In [30]: explainer = shap.Explainer(dtree, X_test)
shap_values = explainer(X_test)
shap.plots.waterfall(shap_values[0])
```



```
In [31]: imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": dtree.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (Decision Tree Regressor)', fontsize=18)
plt.xlabel ('Importance', fontsize=16)
plt.ylabel ('Feature Name', fontsize=16)
plt.show()
```



## AdaBoost Regressor

```
In [33]: from sklearn.ensemble import AdaBoostRegressor
from sklearn.model_selection import GridSearchCV

# Create an AdaBoost Regressor object
ada = AdaBoostRegressor()

# Define the hyperparameter grid
param_grid = {'n_estimators': [50, 100, 150],
              'learning_rate': [0.01, 0.1, 1]}

# Create a GridSearchCV object
grid = GridSearchCV(ada, param_grid, cv=5, scoring='neg_mean_absolute_error')

# Fit the GridSearchCV object to the training data
grid.fit(X_train, y_train)

# Print the best hyperparameters
print("Best hyperparameters: ", grid.best_params_)

Best hyperparameters: {'learning_rate': 0.01, 'n_estimators': 100}
```

```
In [34]: from sklearn.ensemble import AdaBoostRegressor
ada = AdaBoostRegressor(n_estimators=100, learning_rate=0.01, random_state=0)
ada.fit(X_train, y_train)
```

Out[34]: AdaBoostRegressor(learning\_rate=0.01, n\_estimators=100, random\_state=0)

```
In [35]: from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = ada.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))
```

MAE is 142898.84136857474  
MAPE is 1.7810437954933251  
MSE is 88172321437.5971  
R2 score is 0.26353868846015893  
RMSE score is 296938.24515814247

```
In [36]: imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": ada.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (AdaBoost Regressor)', fontsize=18)
plt.xlabel ('Importance', fontsize=16)
plt.ylabel ('Feature Name', fontsize=16)
plt.show()
```

